

## Crash Recovery

## Three Types of Failure


- ◆ Transaction failure
- ◆ Media Failure (Hard Crash)
- ◆ System Failure (Soft Crash)

## Review: The ACID properties

- ◆ Atomicity: All actions in the transaction happen, or none happens
- ◆ Consistency: Transaction transforms DB from consistent states to consistent states
- ◆ Isolation: Execution of one transaction is isolated from that of other transactions
- ◆ Durability: If a transaction commits, its effects persist
- ◆ The Recovery Manager guarantees Atomicity & Durability

## Motivation

- ◆ Atomicity: Transactions may abort
- ◆ Durability: DBMS may stop running



Desired behavior at system restart:

- ◆ T1, T2 & T3 should be durable.
- ◆ T4 & T5 should be aborted (effects not seen).

## Assumptions

- ◆ Concurrency control is in effect
- ◆ Updates are happening "in place", i.e. data is overwritten on (deleted from) the disk

## Buffer Management

- ◆ Features of buffer:
  - ▲ Volatile. Update in buffer may get lost at system crash
  - ▲ Limited space. At some point, no free buffer frame is available
- ◆ What can a DBMS do?
  - ▲ Update data in buffer frames, force write frames to disk if necessary
  - ▲ If no frame is available, steal frame from a running transaction (need to flush updated frames first)

## Force vs Steal: Options

- ◆ Force every update.
  - ▲ Poor response time
- ◆ No force.
  - ▲ What if transaction commits & DBMS crashes before updates are written to disk?
- ◆ No steal.
  - ▲ Poor throughput
- ◆ Steal.
  - ▲ What if transaction aborts after some of its frames are flushed?

## Force vs Steal: Ideas

	No Steal	Steal
Force	Simple	
No Force		Desired

- ◆ No force. Keep enough info on disk at commit time to prepare to redo committed updates
- ◆ Steal. Keep enough info on disk at flush time to prepare to undo aborted updates

## Logging

- ◆ A log is a file of entries about DB updates
  - ▲ Sequentially recorded, first in buffer, then on stable storage
  - ▲ Contain minimal information required to perform undo & redo operations of transactions
- ◆ Recovery Manager keeps a log of DB operations
  - ▲ Record log during normal operation
  - ▲ Use log to redo/undo transactions at system restart time

## Log Entries

- ◆ <T start>
  - ▲ Each transaction has a unique, system generated id
- ◆ <T, X, oldVal, newVal> (or <T, X, newVal>)
  - ▲ oldVal is needed for undo, newVal is needed for redo
- ◆ <T commit>
- ◆ <T abort>

## Write-Ahead Logging (WAL)

- ◆ RM must coordinate the writing of log entries with writing of updated buffer frames
- ◆ The Write-Ahead Logging Protocol :
  - ▲ Must force the log entry for an update before the updated frame is written to disk (to guarantee Atomicity)
  - ▲ Must write all log entries for a transaction before it commits (guarantees Durability)

## When to Update Database?

- ◆ Every update must be done in buffer first
- ◆ Every update must be recorded in log
- ◆ When can DB update start?
  - ▲ Immediate Update : As soon as an update is made in buffer & in log
  - ▲ Deferred Update : As soon as transaction is ready to commit
- ◆ When is DB update done?
  - ▲ Whenever updated frame is stolen
  - ▲ Whenever updated frame is forced out

## Redo and Undo

- ◆ Undo
  - ▲ For each updated data, restore its old value
  - ▲ Must be done in reverse order
  - ▲ Must be repeatable
  - ▲ Needed for recovery & (normal) abort
- ◆ Redo
  - ▲ For each updated data, set to its new value
  - ▲ Must be done in forward order
  - ▲ Must be repeatable
  - ▲ Needed for recovery

## Recovery Algorithms

- ◆ Each algorithm describes
  - ▲ Normal operations:  
read, write, commit, abort
  - ▲ Restart operation
- ◆ Types of algorithms:
  - ▲ Undo/redo
  - ▲ No-undo/redo
  - ▲ Undo/no-redo
  - ▲ No-undo/no-redo

## Undo/Redo Algorithm

- ◆ Buffer Management:
  - ▲ Immediate Update
  - ▲ Steal/no-force
- ◆ Normal Operations:
  - ▲ Update in buffer, record in log, frame can be stolen
  - ▲ Commit places a log entry  $\langle T, \text{commit} \rangle$
  - ▲ Abort performs an undo in buffer, then appends a log entry  $\langle T, \text{abort} \rangle$

## Undo/Redo Algorithm (cont.)

- ◆ Restart Operation:
  - ▲ Undo aborted & uncommitted transactions
    - ✦ Read log backwards & make a commit list (CL) (transactions with  $\langle T, \text{commit} \rangle$  in log)
    - ✦ At meantime, undo every update made by transactions not in CL (done in buffer)
  - ▲ Redo committed transactions
    - ✦ Read log forward and redo every update of committed transactions (in buffer)

## Undo/Redo Algorithm: Example

- ◆ Log on stable storage at system crash:

<T1, start><T1, D, 25, 20><T1, commit>  
 <T4, start><T4, B, 10, 15><T2, start>  
 <T2, D, 20, 12><T4, A, 10, 24><T4, commit>  
 <T3, start><T3, A, 24, 30><T2, D, 12, 35>


- ◆ Restart operations:

- ▲ CL = {T4, T1}
- ▲ Undo T2, T3: D=20, A=24
- ▲ Redo T1, T4: D = 20, B = 15, A = 24

## Undo/Redo Algorithm: Discussion

- ◆ Guarantees both Atomicity and durability.
- ◆ Efficient normal operations
- ◆ Higher degree of concurrency may be obtained.
  - ▲ Intermediate result of a transaction is visible to other transactions, after it is written on disk.
- ◆ May cause cascading abort.
  - ▲ Best used with cascadeless schedules


## How Far To Trace Back?



- ◆ How far should the log be traced back?
  - ▲ May have to go back to the beginning
  - ▲ May perform many unnecessary redo/undo

## Checkpoint: The Idea

- ◆ Create checkpoint at which database is up-to-date and consistent.



- ◆ Only redo T7, T8 and undo T9.

## Checkpoint Operations

- ▲ Stop accepting new transaction
- ▲ Suspend running transactions
- ▲ Force write log buffer frame
- ▲ Force write all updated frames
- ▲ Append <checkpoint, AL> to log, where AL is a list of active transactions at checkpoint time
- ▲ Resume normal execution

## Recovery With Checkpoint

### ◆ Log entries at system crash:

<T1, start><T1, D, 25, 20><T1, commit>  
 <T4, start><T4, B, 10, 15><T2, start>  
 <T2, D, 20, 12><checkpoint, (T2, T4)>  
 <T4, A, 10, 24><T4, commit><T3, start>  
 <T3, A, 24, 30><T2, D, 12, 35>

### ◆ Restart operations:

- ▲ Undo T2, T3 (all operations)
- ▲ Redo T4 (only operations after checkpoint)


## No-Undo/Redo Algorithm

- ◆ Buffer Management:
  - ▲ Deferred Update
  - ▲ No-force, No steal
- ◆ Normal Operations:
  - ▲ Updates are done in buffer & recorded in log, no steal of buffer frames is allowed
  - ▲ Abort appends <T, abort>
  - ▲ Commit appends <T, commit>, and allows steal of buffer frames

## No-Undo/Redo Algorithm (cont.)

### ◆ Restart operation

- ▲ Redo transactions that commit after the most recent checkpoint (in forward order)
- ▲ Restart transactions that are active at the time of crash



### ◆ Redo T2

## No-Undo/Redo: Discussion

- ◆ Guarantees Atomicity and durability
- ◆ Efficient normal operations
- ◆ No undo
- ◆ Lower degree of concurrency since no updated value of a transaction is visible to other transactions before the transaction commits.
  - ▲ All transactions that need to access the result of T have to wait until T completes.
  - ▲ System performance may suffer.